

## 1. Transformer Model

**Dataset Analysis and Preprocessing** The text dataset exhibited significant variance in length distribution:

- Word count range: 13 to 54 words for most entries
- Mean word count: approximately 50 words
- Maximum length: 783 words
- Vocabulary size: approximately 32,000 unique words
- Distribution showed considerable skewness due to outlier entries with exceptionally high word counts
- Standard deviation was notably large, indicating substantial variation in text lengths

### Preprocessing Steps:

1. Text normalization:
  - Conversion to lowercase
  - Removal of non-alphabetic characters
  - Reduction of multiple spaces
2. Linguistic processing:
  - Lemmatization using spaCy
  - Stop word removal
3. Outlier handling:
  - IQR-based outlier removal

### Post-preprocessing Statistics:

- Typical word count range: 7 to 28 words
- Mean word count: approximately 17.36 words
- Standard deviation: 14.94
- Maximum length reduced to 391 words
- 25th-75th percentile range: 6-23 words

### Vectorization and Model Configuration Vectorization Trials:

- Multiple combinations of token limits and vector dimensions were tested.
- Optimal configuration: 5,000 tokens with 30-dimensional vectors.

### Model Details:

- Vocabulary size: 5,000 tokens
- Sequence length: 30 tokens
- Embedding dimension: 100
- Two transformer encoder blocks:
  - Intermediate dimension: 256
  - Number of attention heads: 4
  - Dropout rate: 0.4
- Dense layers:
  - First dense layer: 128 units with ReLU activation
  - Second dense layer: 64 units with ReLU activation
  - Output layer: 5 units with softmax activation
- Regularization:
  - L2 regularization (0.01)
  - Dropout rate: 0.5

### **Training Configuration**

- Optimizer: AdamW with 1e-4 learning rate
- Loss function: Sparse categorical crossentropy
- Batch size: 128
- Early stopping with 5 epochs patience
- Learning rate reduction on plateau

### **Performance**

- Final validation accuracy: 67%
  - Training process revealed initial overfitting.
  - Regularization and learning rate adjustments improved generalization.
- 

## **2. LSTM Architecture**

### **Dataset Analysis and Preprocessing   Preprocessing Steps:**

1. Data scaling:
  - Features were scaled using MinMaxScaler or StandardScaler to normalize data for better convergence.

- Example:

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

## 2. Sequence preparation:

- Data was reshaped into sequences (3D tensors: (samples, time\_steps, features)) to fit RNN input requirements. Example:

```
import numpy as np
time_steps = 50
X_train_seq = []
y_train_seq = []

for i in range(len(X_train_scaled) - time_steps):
    X_train_seq.append(X_train_scaled[i:i + time_steps])
    y_train_seq.append(y_train_scaled[i + time_steps])

X_train_seq = np.array(X_train_seq)
y_train_seq = np.array(y_train_seq)
```

## 3. Label encoding (if classification):

- Labels were one-hot encoded for classification tasks. Example:

```
from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder()
y_train_encoded = encoder.fit_transform(y_train.reshape(-1, 1)).toarray()
```

## 4. Batch preparation (optional):

- Data generators were used for efficient batch preparation. Example:

```
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator
generator = TimeseriesGenerator(X_train, y_train, length=50, batch_size=32)
```

## Model Details:

- Embedding layer: 100-dimensional embedding
- LSTM layer: 128 units
- Dense layers:
  - First dense layer: 64 units with ReLU activation
  - Output layer: 5 units with softmax activation
- Regularization:
  - L2 regularization (0.01)

- Dropout rate: 0.5

### **Training Configuration**

- Optimizer: Adam with 1e-3 learning rate
- Loss function: Sparse categorical crossentropy
- Batch size: 64
- Early stopping with 5 epochs patience

### **Performance**

- Final validation accuracy: 64.6%
  - Demonstrated good handling of sequential patterns.
  - Balanced performance between training and validation.
- 

## **3. GRU Architecture**

### **Dataset Analysis and Preprocessing   Preprocessing Steps:**

#### **1. Text Preprocessing:**

- Removed URLs, mentions, hashtags, and non-alphabetic characters.
- Tokenized the text using `nltk.word_tokenize`.
- Removed stopwords using NLTK's English stopwords list.
- Lemmatized words using `WordNetLemmatizer`.

#### **2. TF-IDF Vectorization:**

- Converted text into numerical representation using `TfidfVectorizer` with a maximum of 5000 features.

3. Label encoding (if classification):

- Labels were one-hot encoded for classification tasks. Example:

```
from sklearn.preprocessing import OneHotEncoder
encoder = OneHotEncoder()
y_train_encoded = encoder.fit_transform(y_train.reshape(-1, 1)).toarray()
```

4. **Data Splitting:**

- Split the dataset into training and testing sets using an 80/20 split.
- Reshaped the data for compatibility with GRU input requirements.

**Model Details:**

- Dense layers:
  - First dense layer: 64 units with ReLU activation
  - Output layer: 5 units with softmax activation
- Regularization:
  - L2 regularization (0.01)
  - Dropout rate: 0.6

**Training Configuration**

- Optimizer: Adam with 1e-3 learning rate
- Loss function: Sparse categorical crossentropy
- Batch size: 64
- Early stopping with 9 epochs patience

**Performance**

- Final validation accuracy: 66.7%
  - Faster training compared to LSTM.
  - Minimal performance trade-off for reduced complexity.
-

### Performance Summary

Model	Validation Accuracy	Key Strengths
Transformer	67.0%	Best contextual understanding
LSTM	64.6%	Strong sequential pattern recognition
GRU	66.7%	Efficient training and simplicity